

2]synthetic dataset

```
from numpy import where
```

```
from collections import Counter
```

```
from sklearn.datasets import make_classification
```

```
from matplotlib import pyplot
```

```
X, y = make_classification(n_samples=1000, n_features=4, n_informative=2,  
                           n_redundant=0, n_classes=3, n_clusters_per_class=1, weights=None,  
                           random_state=1)
```

```
print(X.shape, y.shape)
```

```
counter = Counter(y)
```

```
print(counter)
```

```
for i in range(10):
```

```
    print(X[i], y[i])
```

```
for label, _ in counter.items():
```

```
    row_ix = where(y == label)[0]
```

```
    pyplot.scatter(X[row_ix, 0], X[row_ix, 1], label=str(label))
```

```
pyplot.legend()
```

```
pyplot.show()
```

```

3] find s

import pandas as pd

import io

import numpy as np

uploaded = files.upload()

df2 = pd.read_csv(io.BytesIO(uploaded['ws copy.csv']), header=None)

print("\nThe Given Training Dataset:\n")

print(df2)

print("\nThe most general hypothesis: ['?', '?', '?', '?', '?', '?']\n")

print("\nThe most specific hypothesis: ['0', '0', '0', '0', '0', '0']\n")

X = np.array(df2.iloc[:, :-1])

y = np.array(df2.iloc[:, -1])

print("\nX (features):\n", X)

print("\ny (target):\n", y)

m, n = X.shape

print("\nShape of X (m, n):", m, n)

print("\nThe initial value of hypothesis:")

hypothesis = ['0'] * (n - 1)

print(hypothesis)

print("\nFind S: Finding a Maximally Specific Hypothesis\n")

for i in range(m):

    if y[i] == 'Yes':

        for j in range(n - 1):

            if X[i][j] != hypothesis[j]:

                if hypothesis[j] == '0':

                    hypothesis = list(X[i, :n - 1])

                else:

                    hypothesis[j] = '?'

        print("The hypothesis {0} for training set {0}: ".format(i + 1), hypothesis)

print("\nThe Maximally Specific Hypothesis for the given Training Examples:\n")

print(hypothesis)

```

4] candidate elimination

```
import pandas as pd
```

```
import numpy as np
```

```
df2 = pd.read_csv('/content/ws.csv', header=None)
```

```
print(df2)
```

```
X = np.array(df2.iloc[:, :-1])
```

```
y = np.array(df2.iloc[:, -1])
```

```
print("\nX:")
```

```
print(X)
```

```
print("\ny:")
```

```
print(y)
```

```
m, n = X.shape
```

```
print("\nShape of X (m, n):", m, n)
```

```
print("\nMost specific hypothesis: ['0'] *", n)
```

```
print("Most general hypothesis: ['?'] *", n)
```

```
def candidate_elimination(X, y):
```

```
    print("\nInitialization of specific_h and general_h")
```

```
    specific_h = ['0'] * n
```

```
    print("Specific_h:", specific_h)
```

```
    general_h = ['?'] * n
```

```
    print("General_h:", general_h, "\n")
```

```
    for i in range(m):
```

```
        if y[i] == "Yes":
```

```
            for j in range(n):
```

```
                if specific_h[j] != X[i][j]:
```

```
                    if specific_h[j] == '0':
```

```
                        specific_h[j] = X[i][j]
```

```
                    else:
```

```
                        specific_h[j] = '?'
```

```
            print("Specific_h updated:", specific_h)
```

```
        if y[i] == "No":
```

```
            print("-ve training example:", X[i], "\n")
```

```

new_general_h = []
for gen_h in general_h:
    for j in range(n):
        if X[i][j] != specific_h[j] and specific_h[j] != '?':
            if gen_h[j] == '?':
                new_gen = list(gen_h)
                new_gen[j] = specific_h[j]
                new_general_h.append(new_gen)
            else:
                new_general_h.append(gen_h)
    general_h = new_general_h
    print("General_h updated:", general_h)

return specific_h, general_h

s_final, g_final = candidate_elimination(X, y)

print("\nFinal Specific_h:")

print(s_final)

print("\nFinal General_h:")

print(g_final)

```

5] ID3

```
import pandas as pd

import numpy as np

dataset = pd.read_csv("/content/drive/MyDrive/id3dataset.csv - Sheet1.csv",
                      names=['age', 'income', 'student', 'credit_rating', 'buys_computer'])

def entropy(target_col):
    elements, counts = np.unique(target_col, return_counts=True)

    entropy = np.sum([(-counts[i]/np.sum(counts)) * np.log2(counts[i]/np.sum(counts))
                      for i in range(len(elements))])

    return entropy

def InfoGain(data, split_attribute_name, target_name="buys_computer"):
    total_entropy = entropy(data[target_name])

    vals, counts = np.unique(data[split_attribute_name], return_counts=True)

    Weighted_Entropy = np.sum([(counts[i]/np.sum(counts)) *
                                entropy(data.where(data[split_attribute_name] == vals[i]).dropna()[target_name])
                                for i in range(len(vals))])

    Information_Gain = total_entropy - Weighted_Entropy

    return Information_Gain

def ID3(data, originaldata, features, target_attribute_name="buys_computer", parent_node_class=None):
    if len(np.unique(data[target_attribute_name])) <= 1:
        return np.unique(data[target_attribute_name])[0]

    elif len(data) == 0:
        return

    np.unique(originaldata[target_attribute_name])[np.argmax(np.unique(originaldata[target_attribute_name]
    ], return_counts=True)[1])]

    elif len(features) == 0:
        return parent_node_class

    else:
        parent_node_class =
np.unique(data[target_attribute_name])[np.argmax(np.unique(data[target_attribute_name],
return_counts=True)[1])]

        item_values = [InfoGain(data, feature, target_attribute_name) for feature in features]

        best_feature_index = np.argmax(item_values)
```

```
best_feature = features[best_feature_index]
tree = {best_feature: {}}
features = [i for i in features if i != best_feature]
for value in np.unique(data[best_feature]):
    sub_data = data.where(data[best_feature] == value).dropna()
    subtree = ID3(sub_data, dataset, features, target_attribute_name, parent_node_class)
    tree[best_feature][value] = subtree
return tree

tree = ID3(dataset, dataset, dataset.columns[:-1])
print('\nDisplay Tree:\n', tree)
```

6] ANN ARTIFICIAL NN

```
import numpy as np
```

```
def sigmoid(x):
```

```
    return 1 / (1 + np.exp(-x))
```

```
def sigmoid_derivative(x):
```

```
    return x * (1 - x)
```

```
inputs = np.array([[0, 0], [0, 1], [1, 0], [1, 1]])
```

```
expected_output = np.array([[0], [1], [1], [0]])
```

```
epochs = 10000
```

```
lr = 0.1
```

```
inputLayerNeurons, hiddenLayerNeurons, outputLayerNeurons = 2, 2, 1
```

```
hidden_weights = np.random.uniform(size=(inputLayerNeurons, hiddenLayerNeurons))
```

```
hidden_bias = np.random.uniform(size=(1, hiddenLayerNeurons))
```

```
output_weights = np.random.uniform(size=(hiddenLayerNeurons, outputLayerNeurons))
```

```
output_bias = np.random.uniform(size=(1, outputLayerNeurons))
```

```
print("Initial hidden weights:")
```

```
print(hidden_weights)
```

```
print("Initial hidden biases:")
```

```
print(hidden_bias)
```

```
print("Initial output weights:")
```

```
print(output_weights)
```

```
print("Initial output biases:")
```

```
print(output_bias)
```

```
for _ in range(epochs):
```

```
    hidden_layer_activation = np.dot(inputs, hidden_weights)
```

```
    hidden_layer_activation += hidden_bias
```

```
    hidden_layer_output = sigmoid(hidden_layer_activation)
```

```
    output_layer_activation = np.dot(hidden_layer_output, output_weights)
```

```
    output_layer_activation += output_bias
```

```
    predicted_output = sigmoid(output_layer_activation)
```

```
    error = expected_output - predicted_output
```

```
    d_predicted_output = error * sigmoid_derivative(predicted_output)
```

```
error_hidden_layer = d_predicted_output.dot(output_weights.T)
d_hidden_layer = error_hidden_layer * sigmoid_derivative(hidden_layer_output)
output_weights += hidden_layer_output.T.dot(d_predicted_output) * lr
output_bias += np.sum(d_predicted_output, axis=0, keepdims=True) * lr
hidden_weights += inputs.T.dot(d_hidden_layer) * lr
hidden_bias += np.sum(d_hidden_layer, axis=0, keepdims=True) * lr
print("\nFinal hidden weights:")
print(hidden_weights)
print("Final hidden bias:")
print(hidden_bias)
print("Final output weights:")
print(output_weights)
print("Final output bias:")
print(output_bias)
print("\nOutput from neural network after 10,000 epochs:")
print(predicted_output)
```


7] KNN

```
from sklearn.datasets import load_iris
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import confusion_matrix, classification_report, accuracy_score

iris = load_iris()
print("Feature Names:", iris.feature_names)
print("Iris Data:", iris.data)
print("Target Names:", iris.target_names)
print("Target:", iris.target)

X_train, X_test, y_train, y_test = train_test_split(
    iris.data, iris.target, test_size=0.25, random_state=42)

clf = KNeighborsClassifier()
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:')
print(cm)

print('\nClassification Report:')
print(classification_report(y_test, y_pred))

accuracy = accuracy_score(y_test, y_pred)
print('\nAccuracy:', accuracy)
print('Misclassification Rate:', 1 - accuracy)
print("\nPredicted Data:")
print(clf.predict(X_test))
print("Actual Test Data Labels:")
print(y_test)

diff = y_pred - y_test
print("\nResult of Misclassifications:")
print(diff)

print('Total number of misclassified samples =', sum(abs(diff)))
```

8] NAÏVE BAYES

```
import pandas as pd

from sklearn import tree

from sklearn.preprocessing import LabelEncoder

from sklearn.naive_bayes import GaussianNB

data = pd.read_csv("ws.csv")

print("the first 5 values of data is:\n",data.head())

X=data.iloc[:, :-1]

print("\n the first 5 values of train data is\n",X.head())

y=data.iloc[:, -1]

print("\n the first 5 values of train output is \n",y.head())

le_outlook=LabelEncoder()

X.Outlook = le_outlook.fit_transform(X.Outlook)

le_Temperature = LabelEncoder()

X.Temperature = le_Temperature.fit_transform(X.Temperature)

le_Humidity = LabelEncoder()

X.Humidity = le_Humidity.fit_transform(X.Humidity)

le_Windy = LabelEncoder()

X.Windy = le_Windy.fit_transform(X.Windy)

print("\n now the train data is:\n",X.head())

le_PlayTennis = LabelEncoder()

y = le_PlayTennis.fit_transform(y)

print("\n now the train output is \n",y)

from sklearn.model_selection import train_test_split

from sklearn.naive_bayes import GaussianNB

from sklearn.metrics import accuracy_score

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20)

classifier = GaussianNB()

classifier.fit(X_train, y_train)

print("Accuracy is:", accuracy_score(classifier.predict(X_test), y_test))
```

9] REGRESSION

```
from math import ceil
```

```
import numpy as np
```

```
from scipy import linalg
```

```
def lowess(x, y, f, iterations):
```

```
    n = len(x)
```

```
    r = int(ceil(f * n))
```

```
    h = [np.sort(np.abs(x - x[i]))[r] for i in range(n)]
```

```
    w = np.clip(np.abs((x[:, None] - x[None, :]) / h), 0.0, 1.0)
```

```
    w = (1 - w ** 3) ** 3
```

```
    yest = np.zeros(n)
```

```
    delta = np.ones(n)
```

```
    for iteration in range(iterations):
```

```
        for i in range(n):
```

```
            weights = delta * w[:, i]
```

```
            b = np.array([np.sum(weights * y), np.sum(weights * y * x)])
```

```
            A = np.array([[np.sum(weights), np.sum(weights * x)], [np.sum(weights * x), np.sum(weights * x * x)]])
```

```
            beta = linalg.solve(A, b)
```

```
            yest[i] = beta[0] + beta[1] * x[i]
```

```
            residuals = y - yest
```

```
            s = np.median(np.abs(residuals))
```

```
            delta = np.clip(residuals / (6.0 * s), -1, 1)
```

```
            delta = (1 - delta ** 2) ** 2
```

```
    return yest
```

```
import math
```

```
n = 100
```

```
x = np.linspace(0, 2 * math.pi, n)
```

```
y = np.sin(x) + 0.3 * np.random.randn(n)
```

```
f = 0.25
```

```
iterations = 3
```

```
yest = lowess(x, y, f, iterations)
```

```
import matplotlib.pyplot as plt
plt.plot(x, y, "r.")
plt.plot(x, yest, "b-")
```

10] EM ALG

```
import matplotlib.pyplot as plt
from sklearn import datasets
from sklearn.cluster import KMeans
import pandas as pd
import numpy as np
from sklearn import preprocessing
from sklearn.mixture import GaussianMixture

iris = datasets.load_iris()

X = pd.DataFrame(iris.data, columns=['Sepal_Length', 'Sepal_Width', 'Petal_Length', 'Petal_Width'])
y = pd.DataFrame(iris.target, columns=['Targets'])

model = KMeans(n_clusters=3)

model.fit(X)

plt.figure(figsize=(14, 7))

colormap = np.array(['red', 'lime', 'black'])

plt.subplot(1, 3, 1)

plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[y.Targets], s=40)

plt.title('Real Clusters')

plt.xlabel('Petal Length')

plt.ylabel('Petal Width')

plt.subplot(1, 3, 2)

plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[model.labels_], s=40)

plt.title('K-Means Clustering')

plt.xlabel('Petal Length')

plt.ylabel('Petal Width')

scaler = preprocessing.StandardScaler()

scaler.fit(X)

xsa = scaler.transform(X)

xs = pd.DataFrame(xsa, columns=X.columns)

gmm = GaussianMixture(n_components=3)

gmm.fit(xs)
```

```
plt.subplot(1, 3, 3)

plt.scatter(X.Petal_Length, X.Petal_Width, c=colormap[gmm.predict(xs)], s=40)

plt.title('GMM Clustering')

plt.xlabel('Petal Length')

plt.ylabel('Petal Width')

print('Observation: The GMM using EM algorithm based clustering matched the true labels more closely
than the Kmeans.')
```